AI sycophancy and Its Effects on Civic Formation and Brazilian Democratic Judgement

Jovem Brasil – Pedro Raffaini

Peer reviewed by Pietro Lampronti

1. Introduction

Prompting artificial intelligence is an increasingly important, and decisive, skill in the new era of digital literacy and optimisation. AI usage enables access to information more easily than before, which has an impact, particularly a political one. As this brief will show, when a student asks an AI model already familiar with their work and previous conversations, it tailors its response to familiar concepts, returning a fluent and well-cited argument. Often, this happens to sharpen the view they walked in with in the first place. When prompted about, for example, a change in the principles of sovereignty, the model is likely to return reasoning built upon theories and concepts, such as Weberian theory or Westphalian principles, that the user has engaged with before. Critically, the user may read this as clarification, supported even further through newly-constructed reasoning. Structurally, systems are built to be well-rated by their users, since companies seek satisfied clients. In LLMs, a system built to be well-rated by its user reflects that person's prior conceptions back in a new, well-constructed resolution: because these systems are optimised against preference models that reward agreement over accuracy (Sharma et al., 2023), users frequently receive a more sophisticated version of what they already believed, rather than an independent check on it. This brief explores how such sycophancy, AI's tendency to align its outputs with what it infers a user already believes, *can* shape conditions of belief formation in Brazilian civic life, and proposes policy responses such as mandatory disclosure of opposing viewpoints and pairing classroom AI tools with a human deliberator (see Section 5). It further argues that AI's mode of opinion-influencing is structurally different from prior mechanisms such as filter bubbles and echo chambers, which makes it harder to counteract using existing remedies. The brief does not claim this always happens or that it must, but aims to understand an existing tendency that is measurable and visible in how Brazilians are addressed by automated systems before elections.

Before exploring the impact of AI agreeability in civic learning, and its impact on Brazilian society, this brief argues in favour of AI use in education. The contents of this paper will provide insight and recommend forms to minimise the negative consequences of use of AI in student's lives. For young students that wish to engage in political debate, an AI-as-deliberator is a powerful connection to a politically engaged environment. In other cases where quality of teaching is low, an AI tutor may be the only substantive means of engagement available.

Therefore, the brief assumes AI remains in use in classrooms, and does not argue in favour of its removal.

2. Sycophancy and mechanisms

Sycophancy is defined by researchers as a model's tendency to align output with what it infers the user believes (or wants to hear). Sycophancy may come at the cost of accuracy (Sharma et al., 2023). The Cheng et al. (2026) study found that eleven AI assistants consistently validated users' actions 50% more often than humans did. Once a response matches more closely a user's views, there is a higher probability of it being preferred by human raters (Sharma et al., 2023). Because these preference models themselves reward agreeable answers, optimising model outputs against them trades truthfulness for sycophancy. The study also notes that 12% of US teens turn to chatbots for emotional support (Cheng et al., 2026). An analysis of preference data also estimated that matching a user's beliefs raises the probability that a response is chosen as the preferred answer by 6%, which further reinforces this optimisation pressure. Furthermore, simple opinion statements from AI, such as "I believe the answer is X," increased users' agreement with incorrect beliefs by 63% relative to responses without such statements, across seven LLMs (Wang et al., 2025). Sycophantic answer flipping, the process in which a model suddenly changes its position after being re-prompted, was documented in 58.2% of cases in mathematical and medical queries (Fanous et al., 2025). This mechanism is caused by preference-based training methods, which reward a model that is agreeable more often than one that is truthful (Sharma et al., 2023).

The property of sycophancy does not simply result in skewed outputs: it changes how certain and extreme people's beliefs become, and does so measurably. A study with 3,285 participants introduced four political topics and four different AI models (Rathje et al., 2025). Participants consistently preferred to interact with the AI models that they agreed with the most: across the studies, people were roughly nine percentage points more likely to choose to re-engage with a sycophantic chatbot than with a disagreeable one that had challenged their beliefs (Rathje et al., 2025). The conversations had with sycophantic chatbots increased attitude certainty, whereas chatbots which disagreed, decreased attitude and extremity. Moreover, sycophantic chatbots also inflated people's belief that they are "better than average" on desirable traits such as intelligence or empathy. The problem lies in the blind-spot finding of the study, which is the fact that participants viewed the sycophantic chatbots as unbiased, though viewed disagreeable chatbots as highly biased. This means that the mechanism is invisible to the people which experience it, and also why it cannot be solved by telling students to "think more critically". Personalisation over time of LLMs also increases likelihood of sycophancy. In the study, a condensed user profile in the model's memory had the largest impact. Participants also preferred models which replicated their own political beliefs, specifically when the model could accurately infer those beliefs from conversation (MIT News, 2026). These results documented tendencies under controlled, studied conditions. Such tendencies, however, are not entirely deterministic: they are not guaranteed outcomes of any single interaction, which is precisely what makes the policy interventions proposed in Section 5 worth pursuing.

3. Philosophical frameworks and echo chambers

Epistemic bubbles are social structures where other relevant opinions have been excluded, perhaps accidentally. Echo chambers, as described by Nguyen, are structures where relevant voices are actively excluded and discredited (Nguyen, Episteme 2020). As a key distinction: members of epistemic bubbles lack exposure to relevant, symmetric information, whereas members of echo chambers have come to distrust all "irrelevant sources," namely sources that disagree with the beliefs of such members. Inadvertently, exposing a member of epistemic bubble to the missing information may actually reinforce the echo chamber instead of dissolve it, particularly if that information arrives from sources a person distrust a priori. Unlike epistemic bubbles, the danger of echo chambers is that they are deliberately formed; disagreeing voices are actively discredited and excluded. Nguyen details that echo chambers are formed because members have been trained to distrust outside sources in such that no shared evidence is accepted as neutral by any side, hence there is no method left to settle disagreements empirically (Nguyen, 145). Beyond theoretical concepts, the political impact is significant. Political judgement sits near cognitive islands, it is difficult to entirely distinguish truths from misinformation, which is why political debate often ends in muddled conclusions. No empirical experiment can conclusively settle normative questions such as if an institutional reform is just. The absence of a decisive test is where echo chambers can form, even when parties are acting in good faith. Sycophantic AI mirrors the structure of an echo chamber; it actively reconstructs their own view back to them as if they were an independent confirmation. Logical constructs are made, and the user feels as if their prior beliefs were confirmed through balanced reasoning. Simply showing a user who has reassured themselves of a claim “the other side of argument” does not lift hem out of an echo chamber.

A possible objection could overlook the dangers of sycophancy by considering that search engines hand sources that are already existing, written by someone else with their own visible biases and reputations. In the digital age, curated interface is a considerable influencer of public opinion. Decade-old studies support this claim: Pariser had already made this point in The Filter Bubble (Pariser, 2011), and Sunstein extended a similar argument in #Republic (Sunstein, 2017). Hence, LLM sycophancy is no different from search engine’s selective bias. However, the response lies in the distinction between user interactions. Search engines are examples of static viewing while LLMs are interactive. A search engine feed is unresponsive once objections surface. AI LLMs are able to adapt arguments in real time to specific challenges, which leverages their power of persuasion. Furthermore, the authority of reasoning is subtle in LLMs while clearer in filtered feeds. Partisan blogs and media usually endorse specific views which are visible through the author’s reputation, tone or specific visual markers. Arguments made with generative AI are made in a reasonable voice harder to discount, even when LLMs are performing selective bias (Sunstein, 2017). AI’s mode of opinion influencing is stronger structurally when compared to prior mechanisms that have existed in prior filter bubbles. Its remedies are far harder to be attained.

Socially, the impact of AI sycophancy in echo chambers is dual. It both reasserts divergences in what people conclude, while it converges the architecture they use to reason. Hence, AI models operate on different levels. Overall, diverse ranges of opinions narrow down into poles as AI provides reassurance of position (Doshi & Hauser, 2024). On the other hand, the methodologies by which those poles are reached are similar. An analysis found that 22 different LLMs, across different model families, produced less diverse creative output than a group of 102 humans (Doshi & Hauser, 2024). The effect sizes, a standard statistical measure of how large a difference is between groups, were between 1.4 and 2.2, indicating a difference large enough to be considered substantial by conventional research benchmarks. Trends showed this homogenisation widens with more output (Sourati & Dehghani, 2026). This effectively means that creative output is more diverse amongst humans, which is no surprise. Another study with 118 participants from India and the United States carried out several culturally grounded writing tasks. AI suggestions led most Indian participants to adopt Western writing styles, which signifies a linguistic impact of the widespread use of AI (Agarwal, Naaman & Vashistha, 2025). Both within-culture and cross-culture similarity of essays rose when AI suggestions were used, and Indian and American writing styles converged once AI was introduced. The issue is not which culture's language is being homogenised, but that this convergence is itself representative of a broader cultural and logical narrowing that AI's architecture brings about. This literary and logical narrowing, paired with echo chambers, results in users who are more certain of their own conclusions while drawing on a narrower range of reasoning styles and sources to reach them. This introduces a risk that depends on how homogenous the AI model and its training menu are; nonetheless, this risk does not by itself jeopardise the use of AI.

AI's neutrality framing carries out political argumentative work. When a model presents both sides of an argument (framing power), it has already made invisible editorial choices about what the two legitimate sides are and where a reasonable centre sits. This type of argumentation is useful, though harder to contest than open advocacy. Models that argue specific positions can be disagreed with directly, while models that seem to present balanced arguments make it difficult to distinguish truths from well-marketed, flawed reasoning. The following finding reveals several, sometimes conflicting, effects of sycophantic AI. Audits which tested models across instruments, languages and model families have repeatedly found most major LLMs to be "left leaning" when evaluated on standard political-bias questionnaires, though results vary by model (Santurkar et al., 2026). Notwithstanding this, another 2026 study found that standard political-bias audits partly capture sycophantic accommodation to the political leanings the model infers from the auditor, rather than a fixed political stance of the model itself (Törnberg et al., 2026). This means that the literature on AI political bias is inherently entangled with sycophancy. In other terms, it is difficult to separate an LLM's own political tendencies from the user's, which means AI ends up taking a kind of "net-zero" stance (Törnberg et al., 2026). It mostly accommodates any type of 21st-century coherent political stance and does so while providing seemingly balanced arguments. However, the thesis of this brief is not an inherent critique of AI, but an exploration of its effects on users. A balance can be staged en route to an agreement. The model constructs logical pathways which appear to be legitimate, at times made with sound reasoning, and in others well-marketed flawed reasoning through argumentative framing. It ultimately lands on a conclusion that confirms the user. The user receives a confirmation of a prior belief on the feeling of having weighed it fairly.

4. Impact in civic education and Brazilian society

An audit by the Aláfia Lab tested five major LLMs (ChatGPT, Gemini, Grok, Deepseek and Meta AI) against made up voter personas. The model, prompted for an “ideal candidate”, recommended a candidate matched to each profile. Weeks later, a TSE (Brazil's electoral court) rule came into effect to prohibit AI from doing exactly that. These systems were reconstructing what a possible person with a specific profile would plausibly choose as their candidate. This kind of “profiling” also can happen in civic education, though not as clearly.

Within the context of civic education, the use of AI raises the stakes of the mechanisms discussed above. In most academic subjects, re-assured beliefs carry smaller social costs as when compared to civic studies. In civic education, the effect is considerable: polarisation and under-exercised judgement are some of the consequences of sycophantic AI in echo chambers. One may contend that students have always deliberated as imperfect interlocutors, and that textbooks may include a degree of bias, and this has never been treated as disqualifying. However, this brief does show a confirming tendency and civic learning’s goal being pulled in opposite directions. This then means that the overall effect of AI on civic learning depends less on whether its use is legitimate in principle, and more on whether it is used well in practice. The dependency on use conditions are what create a positive outlook; the underlying problem is addressable through policy design. Citizens without access to a politically engaged environment are especially likely to turn to LLMs to understand or question their own beliefs, though anyone might do so. This generates heightened exposure to the risks discussed in Sections 2 and 3, though not as a guaranteed outcome, and it is precisely this population that policy should prioritise in safeguarding the correct use of AI.

Current Brazilian policy lacks sufficient coverage to properly safeguard the use of AI. However, weighing the consequences discussed above against the current political landscape of polarisation, it should prioritise introducing measures. Currently, both (i) the MEC’s Referencial para o Uso e Desenvolvimento Responsáveis de Inteligência Artificial na Educação (March 2026), and (ii) the parecer of the Conselho Nacional de Educação (CNE, Brazil’s National

Education Council) classifying AI applications in education by risk tiers, stand and were approved in preliminary form in May 2026. However, at the time of writing the CNE's parecer is still to go through public consultation and a plenary vote at the CNE, and then homologation by the Ministry of Education. Within the parecer, there are risk-tiers that are introduced regarding the use of AI, and may be used for flagging source discrepancies within the outputs produced by AI models. There are four tiers identified. Low risk AI tools are in the first tier, covering material organisation and accessibility. The other three tiers have AI tools that get progressively more independent. High-risk applications include automated grading and biometric monitoring. The framework explicitly prohibits emotional surveillance or any sort of social scoring for disciplinary purposes. Intelligent tutoring systems are classified as "moderate-risk" and require full transparency in how they process information. For these tutoring systems specifically, the framework mandates compliance with the LGPD (Lei Geral de Proteção de Dados, Brazil's general data protection law) for data protection of minors. All tiers are built to catch data protection harms, which is useful, though there is no mention of any content-shaping specific policy, or even anything specific to civic education. AI civics tutors sit within the same moderate risk tier as a mathematics tutor, and transparency or reasoning tells nothing about sycophancy or the effects of echo chambers, which is what should be addressed. Further regulations about the public use of AI include the TSE's (Electoral tribunal) 23.755/2026 resolution covering regulation of AI include campaign speeches. However, nothing covers civic education specifically.

5. Solutions

Proposed solutions of this brief include:

1. **New risk tier carved out in the CNE ladder**: For AI models that act in civic education, legislation must account for the risks involved. This could be done by treating political and civic education separately from the current four risk tier categories of data-protection and automated decision harms. This would enable for the following solutions to be implemented with larger clarity.
2. **Friction requirement disclosures:** This solution involves the operating mechanisms of response for AI. Any AI tool that is marketed/can be used for civic education must

surface the strongest available opposing case to the specific view inquired about by users, even unprompted. The AI must activate an automatic prompt that commands this action when prompted about any politically controversial issue.

3. **Canon transparency**: the AI model openly discloses the intellectual traditions that a given answer draws from. This creates transparency as to which sources are used for reasoning and output, making visible for eg. Western-default menus that were used.
4. Human-deliberator pairing: For any sort of civic learning which includes the use of AI, the tool must be paired with a professional human deliberator, such as a teacher, to oversee use. This turns the two uses of AI identified earlier in this brief, AI-as-deliberator and AI-as-tutor, into actual design principles.
5. Additions to national curricula: In subjects such as computer science, introduce curricular modules explaining AI sycophancy and algorithmic echo chambers. This should be taught as an epistemic vigilance practice, equipping students with instruments such as Socratic questioning, so that they may develop a meta-critical thinking ability to identify specific moments where AI may be telling them only what they want to hear.

6. Personal address and reflexive coda

This brief was written with the author's own interactions with LLMs. The author went through a specific series of questions that intended to test the extent to which Claude, ChatGPT and Gemini defaulted to familiar frameworks. For example, when Claude was asked about political science and sovereignty, it defaulted to Weber and Habermas. This generated the intellectual material and questioning methods used for the writing of this brief. The author understands that AI sycophancy carries diverse effects that should be addressed properly, particularly in civic learning. Brazilian political landscape is increasingly shaped by informal sources moved through messaging platforms and social networks. As the democratic environment of discussion is directed to the digital world, citizens must understand fully how to use AI as a beneficial tool. Besides the policies and measures cited above, the author believes that most importantly, citizens should be minimally aware of sycophantic AI and echo chambers as concepts themselves.

Cited Works

"CNE Aprova Diretrizes Sobre Uso De IA Na Educação." Poder360, 12 May 2026, www.poder360.com.br/poder-educacao/cne-aprova-diretrizes-sobre-uso-de-ia-na-educacao/.

Cheng, M., et al. "Sycophantic AI Decreases Prosocial Intentions and Promotes Dependence." Science, vol. 391, 2026, eaec8352, www.science.org/doi/10.1126/science.aec8352.

"Conselho Nacional De Educação Cria Regras Para Uso De IA Em Escolas E Classifica Risco." Portal Olavo Dutra, 2026, www.portalolavodutra.com.br/materia/conselho_nacional_de_educacao_cria_regras_para_uso_de_ia_em_escolas_e_classifica_risco.

Doshi, Anil R., and Oliver P. Hauser. "Generative AI Enhances Individual Creativity but Reduces the Collective Diversity of Novel Content." Science Advances, vol. 10, 2024, eadn5290, www.science.org/doi/10.1126/sciadv.adn5290.

"IA Nas Escolas: O Que Muda Com As Novas Diretrizes Do CNE." CartaCapital, 2026, www.cartacapital.com.br/toquetec/ia-nas-escolas-o-que-muda-com-as-novas-diretrizes-do-cne/.

"Measuring Opinion Bias and Sycophancy via LLM-Based Persuasion." arXiv, 2026, arxiv.org/pdf/2604.21564.

Nguyen, C. Thi. "Echo Chambers and Epistemic Bubbles." Episteme, 2020, www.cambridge.org/core/journals/episteme/article/abs/echo-chambers-and-epistemic-bubbles/5D4AC3A808C538E17C50A7C09EC706F0.

Pariser, Eli. The Filter Bubble: What the Internet Is Hiding from You. Penguin, 2011.

"Personalization Features Can Make LLMs More Agreeable." MIT News, 18 Feb. 2026, news.mit.edu/2026/personalization-features-can-make-llms-more-agreeable-0218.

"A Rational Analysis of the Effects of Sycophantic AI." arXiv, 2026, arxiv.org/pdf/2602.14270.

Rathje, Steven, et al. "Sycophantic AI Increases Attitude Extremity and Overconfidence." OSF Preprints, 2025, doi.org/10.31234/osf.io/vmyek.

Santurkar, Shibani, et al. "Whose Opinions Do Language Models Reflect?" 2023.

Sharma, Mrinank, et al. "Towards Understanding Sycophancy in Language Models." arXiv, 2023, arxiv.org/abs/2310.13548.

Sourati, Zhila, and Morteza Dehghani. "The Homogenizing Effect of Large Language Models on Human Expression and Thought." Trends in Cognitive Sciences, 2026.

Sunstein, Cass. #Republic: Divided Democracy in the Age of Social Media. Princeton UP, 2017.

Törnberg, Petter, et al. “Political Bias Audits of LLMs Capture Sycophancy to the Inferred Auditor.” arXiv, 2026, arxiv.org/html/2604.27633.

“Veja O Que Muda Na Sala De Aula Com As Novas Regras Para IA Do Conselho Nacional De Educação.” Correio 24 Horas, 2026, www.correio24horas.com.br/politica/veja-o-que-muda-na-sala-de-aula-com-as-novas-regras-para-ia-do-conselho-nacional-de-educacao-0526.

“What Research Says About ‘AI Sycophancy’.” Tech Policy Press, 2026, www.techpolicy.press/what-research-says-about-ai-sycophancy/.